Reinforcement Learning for Chatbot with Emotions

Luyi Ma
School of Computer Science
Carnegie Mellon University
Pittsburgh, PA 15213
`luyim@andrew.cmu.edu`

Zhefan Zhu
School of Computer Science
Carnegie Mellon University
Pittsburgh, PA 15213
`zhefanz@andrew.cmu.edu`

1 Introduction

In daily conversations, like online chatting, users send messages with emotional interactions instead of simple Q&A pattern. A natural, fluent one-to-one conversation, generally, contains dialogues in which both users share similar emotion levels, such as happy, sadness. Our goal is to model a natural dialogue and the model can figure out users' emotions and response words with proper emotions.

2 Background

Dialogue systems, such as chatbots have become ubiquitous in modern society. A lot of chatbots have been presented so far. One of the latest chatbots is MILABOT, which is based on deep reinforcement learning, and developed by the Montreal Institute for Learning Algorithms (Serban, Sankar et al. 2017). Figure 1 shows their system control flow. In their chatbot, they have several response models to generate response separately to the input sentence, and the system is also given the dialogue history. The responses generated by sub-models are selected and evaluated according to some selection policies, and the best response is returned to the user.
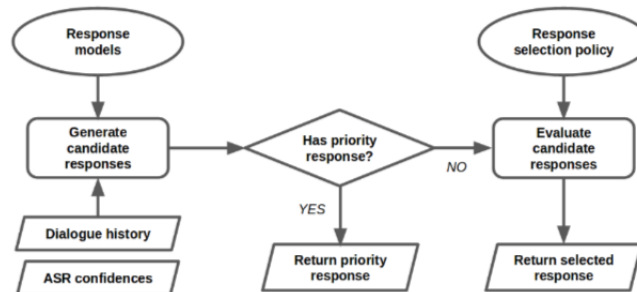

Figure 1: Dialogue manager control flow.

The essential part of a dialogue system is response models. Common response models include knowledge based Q&A systems, template based models, retrieval based neural networks, and generation based neural networks.

For our goal setting, we mainly focus on generation based neural networks. One typical model is the encoder-decoder framework of seq2seq model (Zhou, et al. 2017), where both layers are GRU layers (cho et al. 2014). In this framework, the encoder layer converts input sequence $X = (x_1, x_2, ..., x_n)$ to hidden state $h = (h_1, h_2, ..., h_n)$. And for the decoder layer, it takes in a context vector and previously decoded word to generate its state s_t. The output token is generated by sampling from output probability distribution computed by s_t.

3 Data and Method

3.1 Dataset and Preprocessing

To train the response model, we collected 2.7 million sentences of subtitles from open subtitle database (`http://opus.lingfil.uu.se/OpenSubtitles.php/`). We choose subtitle dataset, because it is difficult to obtain chatting dataset, and subtitles are pretty similar to dalogues at most of the time. The subtitles of movies and TV programs provide sufficient conversation contents for training. Nevertheless, subtitles are slightly different from daily conversation due to theatrical and dramatic contents. We parsed conversations from original data which was in .xml format, and filtered long segment of words which is impractical in daily conversation, which could be monologues in movies.

To train the model determining the emotion of a sentence, we used a dataset of tweets with emotion labels from http://knoesis.org/library/resource.php?id=1749 (Wang et al. 2012). This dataset tags tweets with seven emotions: anger, fear, joy, love, sadness, surprise and thankfulness. The format of the original dataset is: tweetId TAB emotionTag. We implemented a crawler to query tweet content from websites and constructed the training set.

In addition, the word2seq embedding vectors (50 dimensions) are obtained from GloVe (https://nlp.stanford.edu/projects/glove/), which is used to transform words into vectors.

3.2 Long Short Term Memory (LSTM) and Seq2Seq Model

In our task, we use Long Short Term Memory model (LSTM) as a building block to construct more complicated Sequence to Sequence model (Seq2Seq Model). The basic LSTM model is usually used for parsing sequential data. Lets say we have a sequence of inputs $x_1, x_2, ..., x_T \in \mathcal{R}^m$. With an initial cell state $c_0 = \mathbf{0} \in \mathcal{R}^h$ and an initial output $h_0 = \mathbf{0} \in \mathcal{R}^h$, inputs can be parsed by LSTM model with h hidden units. At time $t \in [1, T]$, LSTM model does the following updates:

$$i_t = \sigma(x_t^T W_i + h_{t-1}^T U_i + b_i)$$
$$f_t = \sigma(x_t^T W_f + h_{t-1}^T U_f + b_f)$$
$$o_t = \sigma(x_t^T W_o + h_{t-1}^T U_o + b_o)$$
$$\hat{c}_t = \tanh(x_t^T W_c + h_{t-1}^T U_c + b_c)$$
$$c_t = f_t \odot c_{t-1} + i_t \odot \hat{c}_t$$
$$h_t = o_t \odot \tanh(c_t)$$

, where $\sigma(\cdot)$ is the sigmoid function (element-wise) and $\odot$ is the element-wise product.

For the Seq2Seq recipe in our task, we adapt the multilayer LSTM model from Oriol et al.[Grammar as a Foreign Language]. Figure 2 shows the multilayer LSTM model. $LSTM_{in}^i$ indicates the i-th encoding LSTM layer and $LSTM_{out}^i$ indicates the i-th decoding LSTM layer.


Figure 2

3.3 Emotion in dialogue utterances

Our task is to build a chatbot with the ability of considering the emotion of the utterances and react with a proper emotion. To extract or predict the emotion/mood of utterances, we use an one layer LSTM model to parse sentences and predict the type of emotions. Basically, we take the last hidden state, h_T, of the LSTM model for emotion classification:

$$O = h_T^T W_1 + b_1$$
$$p = SoftMax(O)$$
$$loss = \frac{1}{N_c} \sum (y_i \log(p_i) + (1 - y_i) \log(1 - p_i))$$

, where $SoftMax(\cdot)$ is the softmax function, y_i is the true label and N_c is the number of emotion types.

We use this model to measure the emotion of utterances in the dialogue.

3.4 Reinforcement Learning for Chatbot

During the training of chatbot, it is quite likely to generate dull sentences which may stop the conversation, or sentences with similar meaning of previous sentences which leads to an infinite cycle. In order to model the long-term influence of generated dialogues and emotion flows between utterances, we adapt the reinforcement learning framework from Li et al.[Deep Reinforcement Learning for Dialogue Generation] with some modifications. We will follow the notation used in [Deep Reinforcement Learning for Dialogue Generation] and define the reinforcement learning framework below.

3.4.1 Framework of Reinforcement Learning

In order to model the generation of dialogues, we simulate this process with two bots, Agent P and Agent Q. A dialogue is represented as a sequence of tuple $[p_i, q_i]$ and an action from an agent is denoted as $a_{\cdot i}$ which can be embedded by a Seq2Seq model (h_{p_i} for a_{p_i} and h_{q_i} for a_{q_i}).

The policy model p_{RL} model the decision making. For example, $p_{RL}(p_{i+1}|p_i, q_i)$ is the conditional likelihood of p_{i+1} given prior state $[p_i, q_i]$). The rewards $R(a_i, [p_i, q_i])$ represents the rewards of action a_i given the prior state $[p_i, p_i]$. The rewards can be four-fold. To avoid dull answering, a reward r_1 is added and can be defined as:

$$r_1 = -\frac{1}{N_\mathcal{S}} \sum_{s \in \mathcal{S}} (\frac{1}{N_s} \log p_{seq2seq}(s|a))$$

, where $\mathcal{S}$ is a set of customized dull sentences. The negative log likelihood $p_{seq2seq}(s|a)$ is maximized and normalized by the size of s and $\mathcal{S}$, $N_\mathcal{S}$ and N_s.

To maintain the information flow, a second reward is considered.

$$r_2 = -\log cosSim(h_{p_i}, h_{p_{i+1}})$$
$$= -\log(\frac{h_{p_i} h_{p_{i+1}}}{||h_{p_i}||||h_{p_{i+1}}||})$$

, where $cosSim()$ computes the cosine similarity. This reward is maximized to penalize semantic similarity between two consecutive turns and avoid "dead cycles".

To optimize the semantic coherence, a third reward r_3 is added.

$$r_3 = \frac{1}{N_a} \log p_{seq2seq}(a|p_i, q_i) + \frac{1}{q_i} \log p_{seq2seq}^{backward}(q_1|a)$$

r_3 is maximized to optimize the forward reward and backward reward.

Finally, we add a reward of emotion adequacies to the model.

$$r_4 = e_{a|p_i, q_i} \log e_{q_i}$$

, where $e.$ is the embedding of utterance emotion from the LSTM model we used to predict sentences' emotions.

Now $R(a, [p_i, q_i])$ can be expressed as,

$$R(a, [p_i, q_i]) = \sum \lambda_i r_i$$
$$\sum \lambda_i = 1$$

, where λ_i is the weight of the reward r_i. Initially, we set $\lambda_1 = \lambda_2 = \lambda_4 = 0.2$ and $\lambda_3 = 0.4$.

The objective function of our model which is optimized during training is,

$$\mathcal{J}(\theta) = E_{p_{RL}}[\sum_{i=1}^{T} R(a_i|[p_i, q_i])]$$

4 Experiments & Results

4.1 Tweet classification

We trained the simple LSTM model with 80% of the total tweet dataset, and used remaining 20% for evaluating the model performance. Since every sample tweet has different length of words, the batch size of the model is set to 1, and Figure 3 shows the trend of test classification error vs. number of training samples. The model converges very fast, and achieves minimum testing error after 20000 iterations (tweet samples). Table 1 shows the prediction results of some sentences picked randomly from websites. We can notice that our model tends to give positive predictions to the inputs, which requires further study.
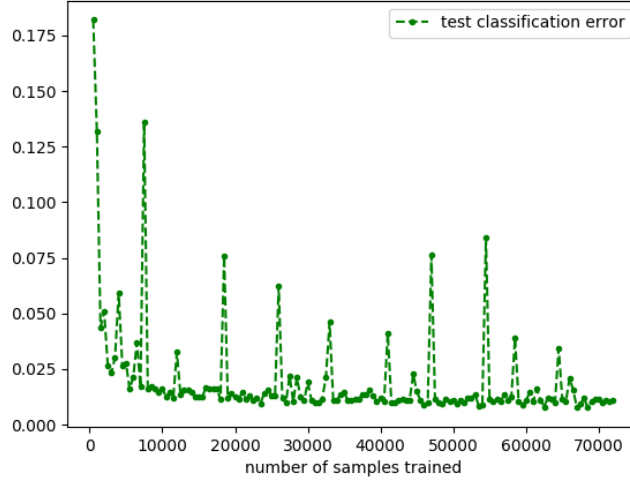

Figure 3: Test classification error of simple LSTM model.

We also try a version of LSTM with a soft attention.

$$\{h_1, h_2, ..., h_T\} = LSTM(\mathbf{X})$$
$$u_i = v^T h_i$$
$$a_i = SoftMax(u_i)$$
$$x_{weighted} = \sum a_i x_i$$

4

, where v is a vector for attention computation. We use the weighted sum of words to predict the type of emotions since usually the emotion is borne in these words with emotion potential like "happy" and "congratulation". Here are some result of classification with attention weights. However, this assumption doesn't work well in the real, complicate world. But it seems to overfit on some words without considering the semantic information of the sentence. So we abandon this model and use a simple LSTM for emotion extraction.

Sentence	Predicted emotion	True emotion
I so pissed . Roger just stabbed me in the back .	joy	anger
I so frustrated .	sadness	sadness
It's so frustrating working with him .	love	anger
I was so frustrated , I stopped caring about the outcome .	fear	anger
I feeling pretty good right now .	joy	joy
I in a very good mood .	fear	joy
It feels so so good taking a long vacation .	joy	joy
I got everything I ever wanted. I feel so blessed .	joy	joy

Table 1: Predictions of some random sentences.

References

[1] Cho, K., Van Merriënboer, B., Bahdanau, D., & Bengio, Y. (2014). On the properties of neural machine translation: Encoder-decoder approaches. arXiv preprint arXiv:1409.1259.

[2] Serban, I. V., Sankar, C., Germain, M., Zhang, S., Lin, Z., Subramanian, S., ... & Mudumba, S. (2017). A Deep Reinforcement Learning Chatbot. arXiv preprint arXiv:1709.02349.

[3] Wang, W., Chen, L., Thirunarayan, K., & Sheth, A. P. (2012, September). Harnessing twitter" big data" for automatic emotion identification. In Privacy, Security, Risk and Trust (PASSAT), 2012 International Conference on and 2012 International Confernece on Social Computing (SocialCom) (pp. 587-592). IEEE.

[4] Zhou, H., Huang, M., Zhang, T., Zhu, X., & Liu, B. (2017). Emotional Chatting Machine: Emotional Conversation Generation with Internal and External Memory. arXiv preprint arXiv:1704.01074.